

Lab3 Design Document

Mei Yixuan 2019011041 Yao92

May 16, 2021

1 Advanced Cache Replacement Policies

1.1 LRU-LIP

In LRU-LIP, we set counter of the newly inserted block as the number of valid ways in the corresponding set minus one. This ensures that newly added blocks are in the least important position. Also, the counters are continuous (i.e. if we have 3 valids ways, their counters have value 0, 1 and 2). This nice property makes eviction and reversion much easier: we can simply use the same function as in LRU. The hardware control overhead of LRU-LIP is one counter each way, which is identical to LRU. Exact hardware cost is in the following figure.




Figure 1: LRU-LIP Hardware Cost

1.2 DIP

In DIP, we use the first set (SET0) as MIP sample and the second set (SET1) as LIP sample. In each memory access on non-dueling sets (i.e. sets other than SET0 and SET1), we update cache control fields according to current better replacement policy. We reset all counters every 100000 memory accesses to avoid potential risk of overflow. Thanks to the good property of LRU-LIP, eviction and reversion of DIP is also identical to LRU. Besides the counter in each way, DIP also needs five counters for data recording. Exact hardware cost is in the following figure.




Figure 2: DIP Hardware Cost

1.3 RRIP

In RRIP, we use counter field of each way to store its RRI. It has field width of 3 bits. Upon hit, we set RRI of corresponding entry as 0. Upon insertion, we set RRI of corresponding entry as long RRI (i.e. 6). When choosing a block for eviction when all ways are occupied, we choose the block with largest RRI and normalize all values to distant RRI (this is identical to adding one repeatedly). Exact hardware cost is in the following figure.




Figure 3: RRIP Hardware Cost

1.4 Results

Results of different replacement policies:

Design	Benchmark 1 Miss Rate	Benchmark 2 Miss Rate	Benchmark 1 Total Cycles	Benchmark 2 Total Cycles
No Cache	100	100	2083191	23508468
Random Replacement	16.47	2.44	1230758	9275269
LRU Replacement	13.86	1.99	1204202	9209763
LRU-LIP Replacement	14.77	2.76	1213470	9322490
DIP-Replacement	14.66	2.27	1212268	9250427
RRIP-Replacement	13.71	2.24	1202611	9246016

Results of different cache capacities and mapping schemes:

Capacity and Mapping	Benchmark 1 Miss Rate	Benchmark 2 Miss Rate	Benchmark 1 Total Cycles	Benchmark 2 Total Cycles
16-entry direct-mapped	29.71	26.23	1365868	12745909
16-entry 4-way associative	28.74	7.52	1356019	10016283
16-entry 8-way associative	29.67	3.56	1365525	9439026
16-entry fully-associative	30.32	3.53	1372080	9433830
32-entry direct-mapped	19.01	19.73	1256732	11798043
32-entry 4-way associative	13.86	1.99	1204202	9209763
32-entry 8-way associative	12.72	1.92	1192565	9199480
32-entry fully-associative	12.67	1.90	1191977	9196845
64-entry direct-mapped	13.80	6.51	1203530	9868691
64-entry 4-way associative	6.49	0.54	1128914	8997426
64-entry 8-way associative	6.24	0.43	1126334	8981639
64-entry fully-associative	6.18	0.39	1125792	8975471

2 Skewed-Associative Cache

2.1 Implementation

In skewed-associative cache, we need separate hash functions for each way. Since we want blocks to be independently and uniformly mapped in each way, we choose a group of universal hash functions for mapping. Specifically, we choose $h_{ab}(k) = (a * k + b) \bmod m$, in which k is the block address, m is the number of sets. In order for this group of hash functions to be universal, we need to choose a and b uniformly random for each way. Considering actual performance, we choose $a = 1$ for all ways, and b be a random number for each way. Note that the tag function used for non-skewed cache is not suitable for skewed cache, since multiple blocks with the same tag may be mapped into the same way of the same set. Therefore, we use the whole address field as tag instead.

2.2 Results

Capacity and Mapping	Benchmark 1 Miss Rate	Benchmark 2 Miss Rate	Benchmark 1 Total Cycles	Benchmark 2 Total Cycles
16-entry 4-way associative	28.74	7.52	1356019	10016283
16-entry 4-way skewed-associative	28.53	8.92	1353896	10220843
32-entry 4-way associative	13.86	1.99	1204202	9209763
32-entry 4-way skewed-associative	13.98	2.02	1205348	9214083

3 Benchmarks

3.1 Replacement Benchmarks

```
.text
# lw a[0], a[32], a[64]
# for i = 0; i < 1000; i++
# lw a[128], a[96]
la t0, argument    # t0: array a
li t1, 1000         # t1: iter
mv t2, zero        # t2: i
lw t3, 0(t0)
lw t3, 32(t0)
lw t3, 64(t0)
loop:
    addi t2, t2, 1
    lw t3, 128(t0)
    lw t3, 96(t0)
    bne t2, t1, loop    # if i != iter then target
```

(a) LRU

```
.text
# for i = 0; i < 1000; i++
# lw a[0], a[32], a[64], a[96], a[128]
la t0, argument    # t0: array a
li t1, 1000        # t1: iter
mv t2, zero        # t2: i
loop:
    addi t2, t2, 1
    lw t3, 0(t0)
    lw t3, 32(t0)
    lw t3, 64(t0)
    lw t3, 96(t0)
    lw t3, 128(t0)
    bne t2, t1, loop    # if i != iter then target
```

(b) LRU-LIP

Figure 4: Replacement Benchmarks

The first program will have better performance on LRU. In this example, LRU will keep `a[96]` and `a[128]` in cache, while LRU-LIP will evict them alternately (note that all five elements are hashed into the same set). Therefore, LRU has very low miss rate while LRU-LIP has 100% miss rate. In the second example, LRU-LIP has better performance. In this example, LRU will cyclically evict blocks from cache, causing all misses. LRU, on the contrary, only incurs two misses each iteration. Therefore, it has much lower miss rate. In actual design of a processor, we may prefer LRU if the programs running on top of it have very good temporal locality, i.e. elements visited will soon be revisited again. If considerable amounts of memory accesses are random access that won't repeat itself, LRU-LIP may be a better choice.

Replacement Policy	bench_lru Miss Rate	bench_lrulip Miss Rate	bench_lru Total Cycles	bench_lrulip Total Cycles
LRU	0.25	100	10055	59006
LRU-LIP	100	40.06	26039	35030

3.2 Writehit Benchmarks

```
.text
# lw a[0], a[32], a[64]
# for i = 0; i < 1000; i++
# sw a[96], a[128]
la t0, argument # t0: array a
li t1, 1000      # t1: iter
mv t2, zero      # t2: i
lw t3, 0(t0)
lw t3, 32(t0)
lw t3, 64(t0)
loop:
    addi t2, t2, 1
    sw zero, 128(t0)
    sw zero, 96(t0)
    bne t2, t1, loop # if i != iter then target
```

Figure 5: Writehit Benchmarks: Write-back

In bench_writehit_back, we repeated write to two blocks in cache. If we use write-back mode, no memory access will actually get to the main memory. However, if we use write-through mode, each and every access will be written back and cause a 10 cycle stall. However, since both modes don't affect cache state, they will have identical miss rate on every program. In this lab's settings, there is no way write-through will have better performance than write-back. This is because write-back only incurs a 10 cycle stall when eviction, while write-through, without help of write buffer, will stall on every access. In actual implementation, write-through may be faster (with write buffer) and less expensive (no dirty bit).

Write Hit Policy	bench_through Miss Rate	bench_back Miss Rate	bench_through Total Cycles	bench_back Total Cycles
Write-back	N/A	0.25	N/A	10055
Write-through	N/A	0.25	N/A	26039

3.3 Writemiss Benchmarks

4 Ripes Bug Report

5 Question Answering